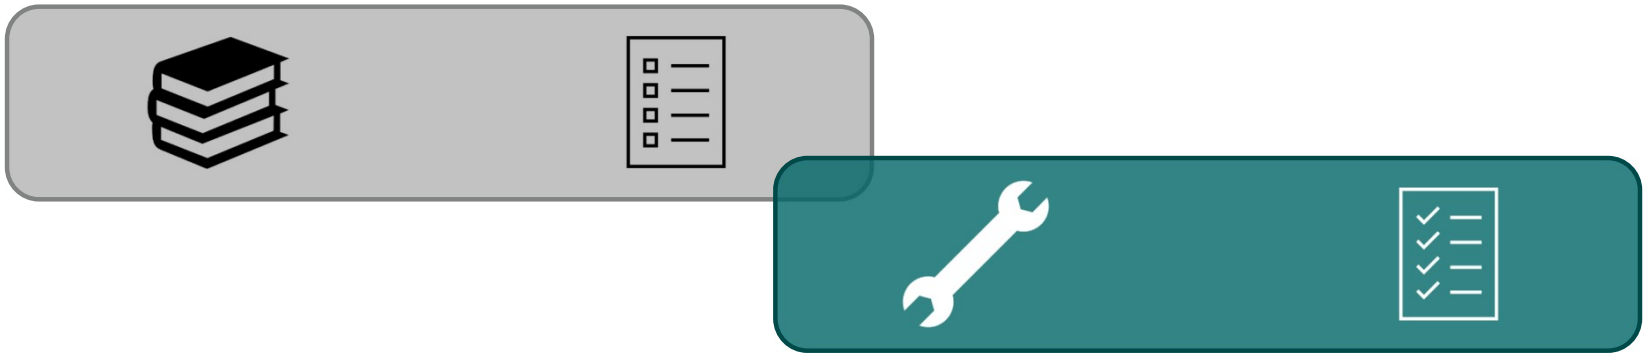

Reproducibility in Data Preprocessing: An Evaluation of Open Source Tools

Simon Grimm

Content

Introduction

Fundamentals

Related Work

Methods and Design

Results

Discussion

Summary and Outlook

Q&A

Introduction

Motivation | Purpose of this Work

- **Reproducibility** - vital and underpins trust in science
  - Ongoing and enhanced focus across different scientific domains and industries [Soh23]
  - Prevalence of data work as a major challenge [Fei+20]

- **Data Preprocessing** – make data suitable for analysis
  - foundation for data mining [AKV19], data science projects [ATSO17], data analysis [Fam+97], machine learning
  - impacts any derived conclusions, model quality, and model fairness [GZ19] [BR21]

- **Open Source** - „The bigger the problem, the more developers are drawn, like magnets, to work on it" [BCG21]
  - Integral to business
  - benefits reproducibility by fostering trust, enabling collaborative work, and emphasizing the value of software and data as artifacts for learning and sharing knowledge [Bar22]

Introduction

Investigated research questions:

RQ1: What are the requirements for reproducible data preprocessing?

RQ2: What are promising open-source tools to enable reproducible data preprocessing?

RQ3: To what extent do existing open-source tools support reproducible data preprocessing?

Fundamentals

- Non-conform standard across different scientific domains [GFI16]
- Often adapted according to a specific context [GK18]
- The Association for Computing Machinery proposes the following terminology[Noab]:
  - **Repeatability**: Same team, same experimental setup
    - *"a researcher can reliably repeat her own computation"*
  - **Reproducibility**: Different team, same experimental setup
    - *"an independent group can obtain the same result using the author's own artifacts"*
  - **Replicability**: Different team, different experimental setup
    - *„independent group can obtain the same result using artifacts which they develop completely independently.*"

Fundamentals

Reproducibility — Data Preprocessing

Data preprocessing comprises all necessary concepts and methods to transform raw data to a high-quality data product that satisfies the requirements for further usage [HKP12].

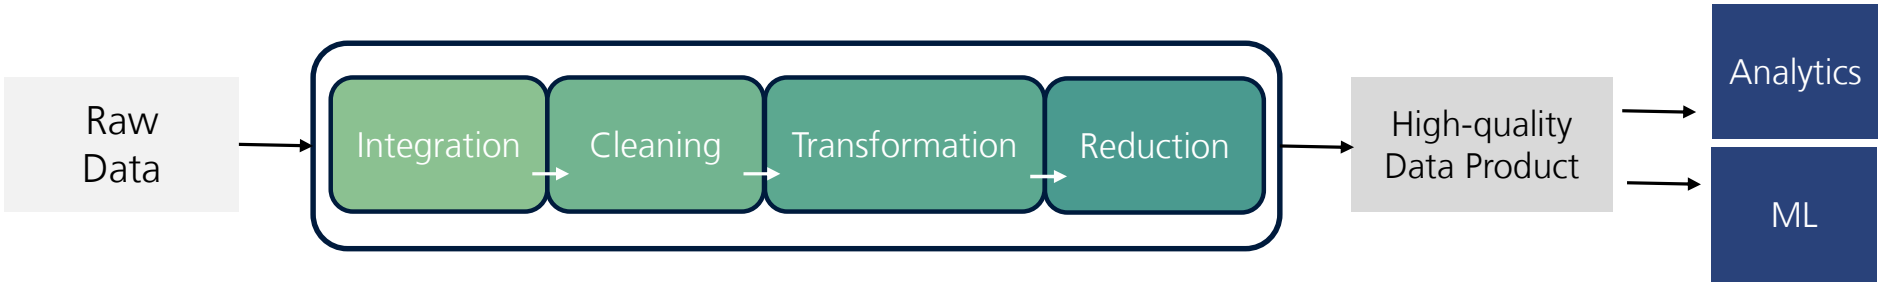

[HKP12]

Fundamentals

Reproducibility — Data Preprocessing

Data preprocessing comprises all necessary concepts and methods to transform raw data to a high-quality data product that satisfies the requirements for further usage [HKP12].

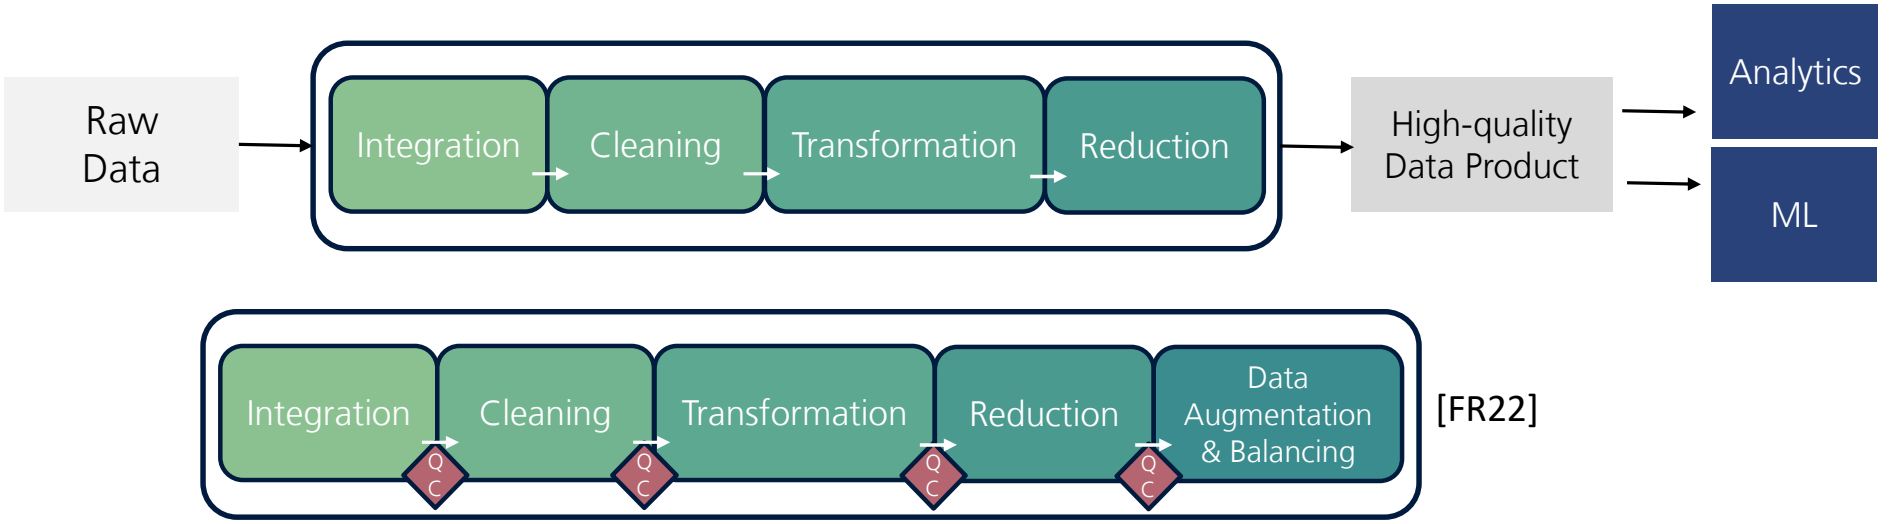

Related Work

Gundersen et al. Albertoni et al.

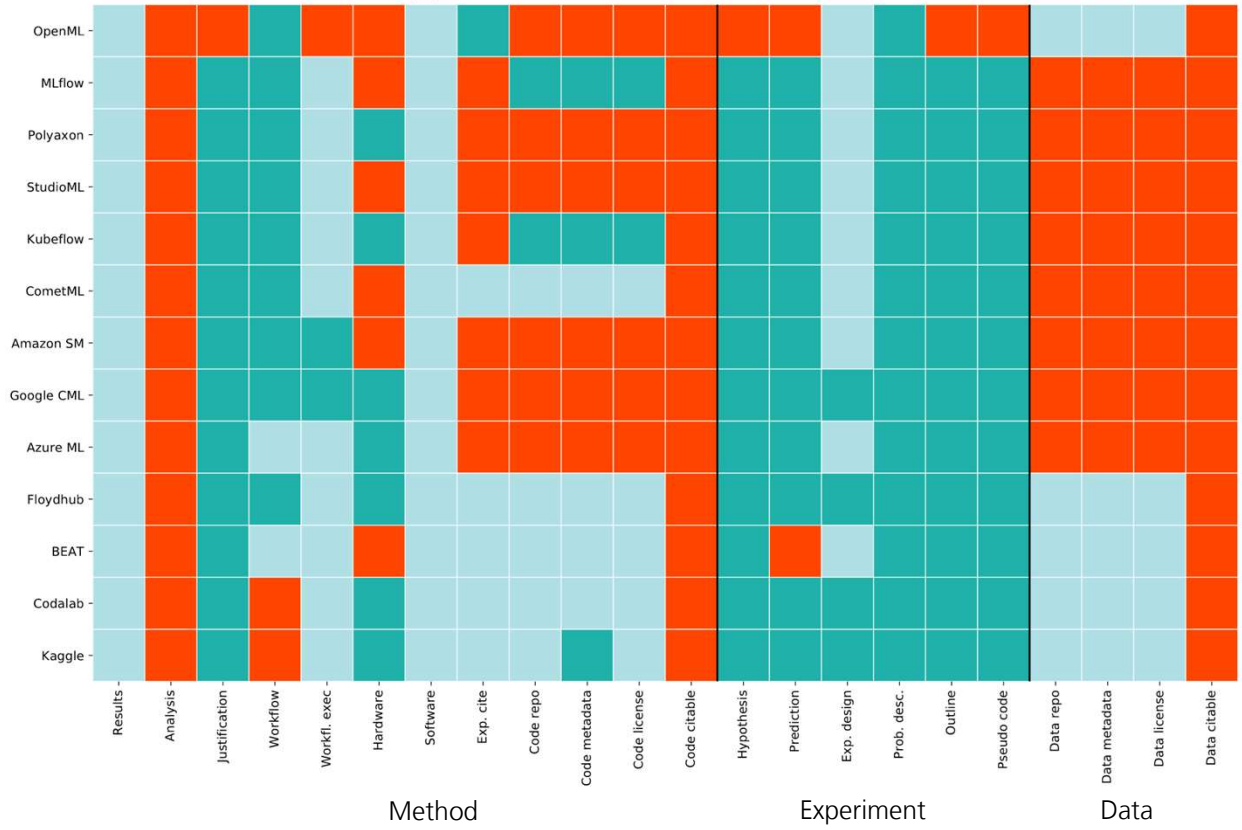

"Do machine learning platforms provide out-of-the-box reproducibility?" [GSI22]

Related Work

Gundersen et al.

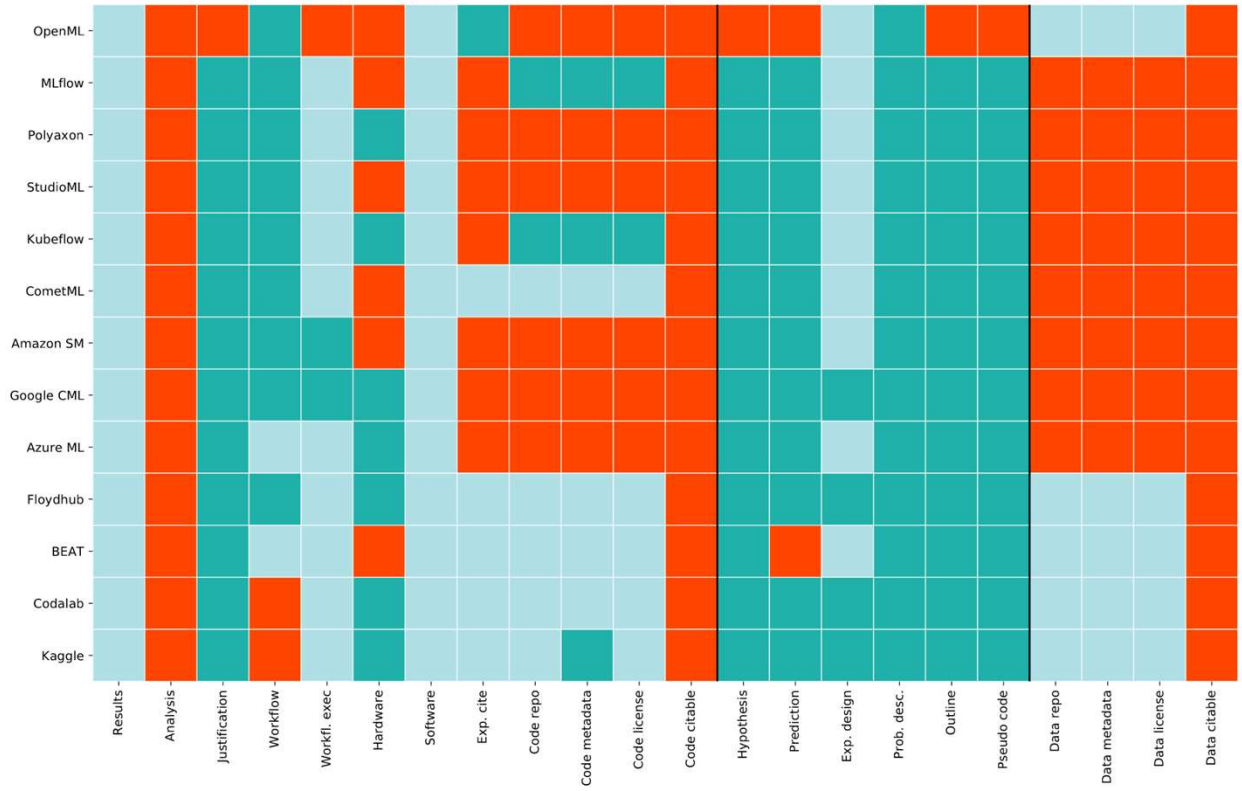

Albertoni et al.

"Do machine learning platforms provide out-of-the-box reproducibility?" [GSI22]

The three reproducibility metrics are defined as follows:

$$R1F(p) = \frac{\delta_1 Method(p) + \delta_2 Data(p) + \delta_3 Exp(p)}{\delta_1 + \delta_2 + \delta_3} \quad (1)$$

$$R2F(p) = \frac{\delta_1 Method(p) + \delta_2 Data(p)}{\delta_1 + \delta_2}, \quad (2)$$

$$R3F(p) = Method(p), \quad (3)$$

Reproducibility metric scores for the 13 platforms.

Platform	R1F	R2F	R3F
OpenML	0.39	0.46	0.17
MLflow	0.33	0.29	**0.58**
Polyaxon	0.32	0.29	**0.58**
StudioML	0.31	0.29	**0.58**
Kubeflow	0.36	0.29	**0.58**
CometML	0.42	0.29	**0.58**
Amazon SM	0.29	0.29	**0.58**
Google CML	0.28	0.25	0.50
Azure ML	0.33	0.29	**0.58**
Floydhub	**0.65**	**0.63**	0.50
BEAT	**0.65**	**0.63**	0.50
Codalab	0.64	**0.63**	0.50
Kaggle	0.63	**0.63**	0.50

Related Work

Gundersen et al.

Feature	Recommendation	Where	Source guideline
Data repository	Share data in a community repository or the simulation enviroment	P	Gundersen et al. [55] Pineau's checklist v2 [114] IJCAI 22 Guideline [140]
Data distribution	How will the dataset will be distributed (e.g., tarball on website, API, GitHub)?	P,M	Datasheets [43]
Data appendix	All novel datasets introduced in this paper are included in a data appendix	S	IJCAI 22 Guideline [140]
Dataset from literature	All datasets drawn from the existing literature (potentially including authors' own previously published work) are publicly available	P,S,M	IJCAI 22 Guideline [140]
Cite Data	All datasets drawn from the existing literature (potentially including authors' own previously published work) are accompanied by appropriate citations	P,S,M	IJCAI 22 Guideline [140]
Data citeable	Generate DOI or PURL.	P,M	Gundersen et al. [55] Datasheets [43]
Data relevant statistic	For all datasets used, The relevant statistics, such as number of examples	P,S,M	Pineau's checklist v2 [114]
Unavailable Dataset Description	All datasets that are not publicly available (especially proprietary datasets) are described in detail	S	IJCAI 22 Guideline [140]
Data collection, annotation and quality	For all datasets used, For new data collected, a complete description of the data collection process, such as instructions to annotators and methods for quality control	P,S,M	Pineau's checklist v2 [114] Datasheets [43]
Train/validation/test splits.	For all datasets used, The details of train/validation/test splits.	P,M	Pineau's checklist v2 [114]
Excluded data	For all datasets used, An explanation of any data that were excluded, and all preprocessing steps.	P,S,M	Pineau's checklist v2 [114]
Preprocessing cleaning and labelling	Was any preprocessing/cleaning/labeling of the data done (e.g., discretization or bucketing, tokenization, part-of-speech tagging, SIFT feature extraction, removal of instances, processing of missing values)? If so, please provide a description.	P,S,M	Datasheets [43]
Raw data	Was the "raw" data saved in addition to the preprocessed/cleaned/labeled data (e.g., to support unanticipated future uses)? If so, please provide a link or other access point to the "raw" data.	P,S,M	Datasheets [43]
Preprocessing software	Is the software used to preprocess/clean/label the instances available? If so, please provide a link or other access point.	P,S,M	Datasheets [43]
Data metadata	Include basic metadata describing the data.	P,M	Gundersen et al. [55]
Dataset contacts	How can the owner/curator/manager of the dataset be contacted (e.g., email address)?	P,M	Datasheets [43]
Data license, Intelectual property, term of use	Give the data a license including Intelectual property and use terms or regulatory restrictions	P,M	Gundersen et al. [55] Datasheets [43]

Table 3. Recommendation for data. The first and second columns summarize what to describe; the third is where the description is likely to be provided (i.e. in metadata (M), the platform (P), or the scientific material, paper, report etc. (S)); the fourth column includes the guidelines from which the recommendation comes

Albertoni et al.

"Reproducibility of Machine Learning:

Terminology, Recommendations and Open Issues" [Alb+23].

Methods and Design

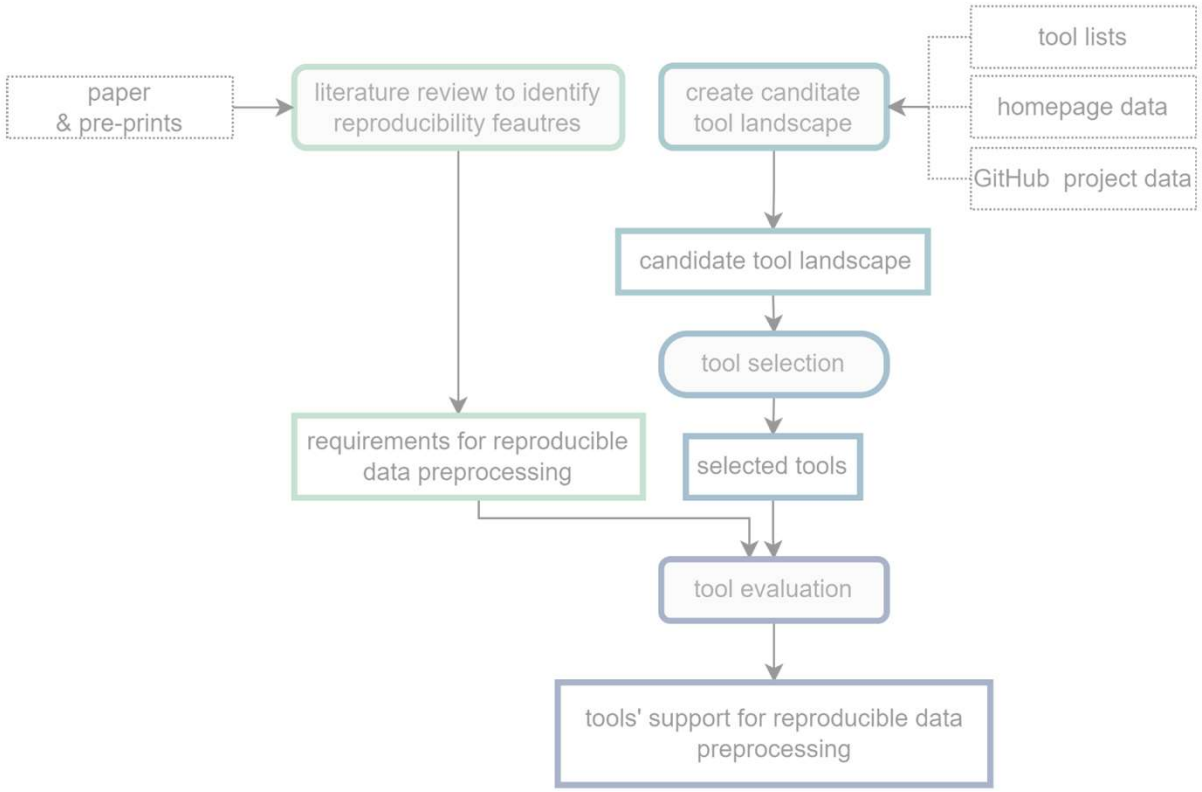

Methods and Design

Overview | Requirements for Reproducibility | Candidate Tool Landscape | Tool Selection | Tool Evaluation

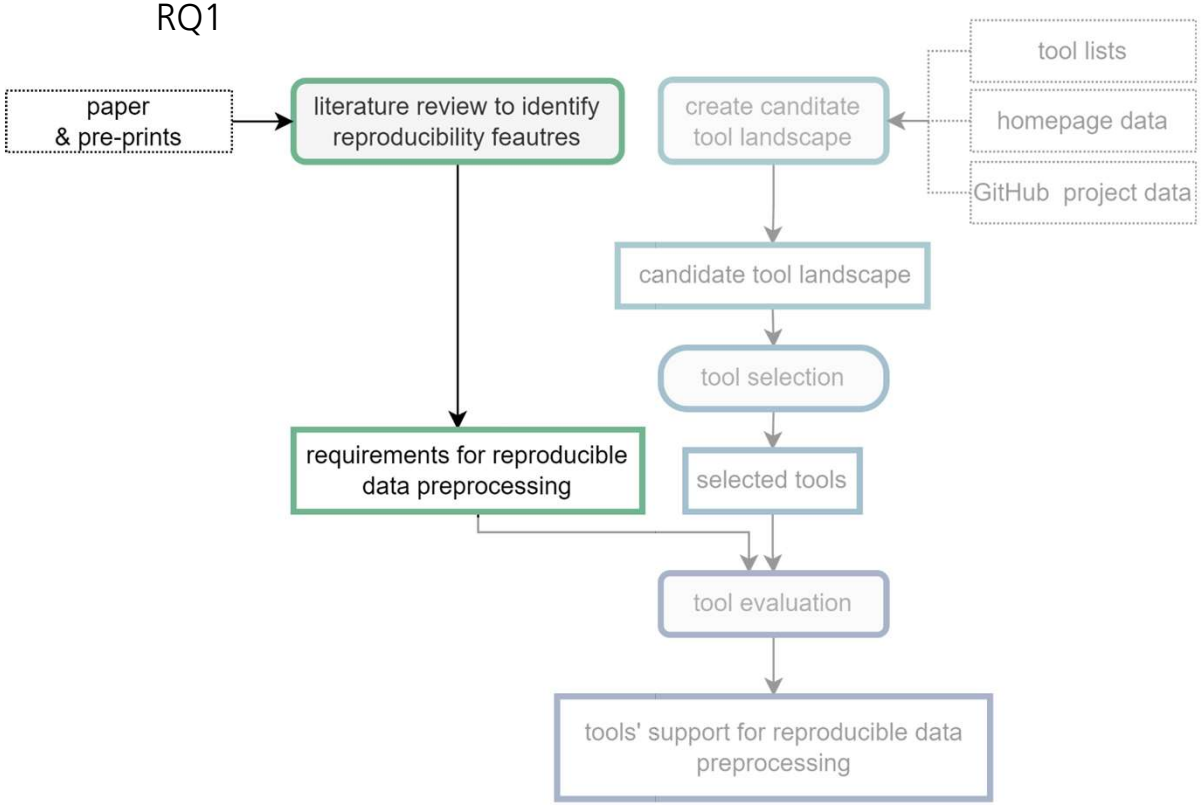

Methods and Design

Overview | Requirements for Reproducibility | Candidate Tool Landscape | Tool Selection | Tool Evaluation

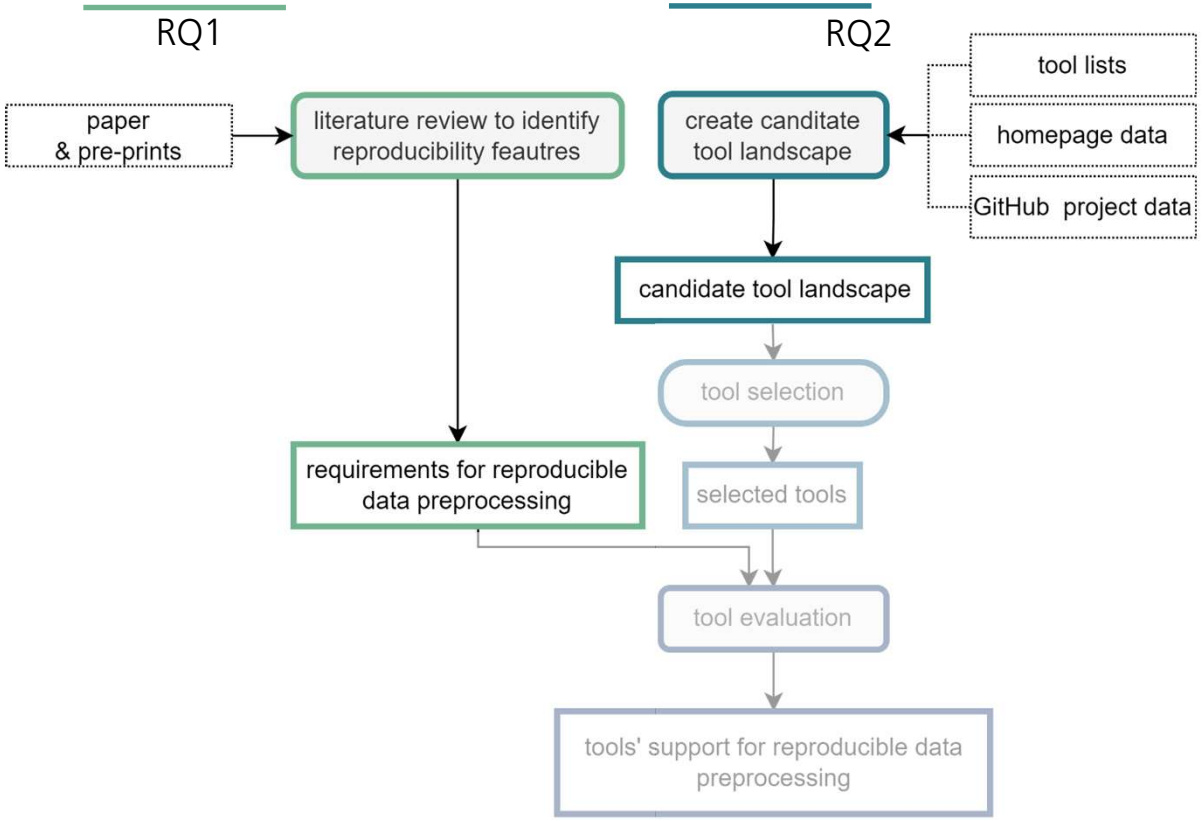

Methods and Design

Overview | Requirements for Reproducibility | Candidate Tool Landscape | Tool Selection | Tool Evaluation

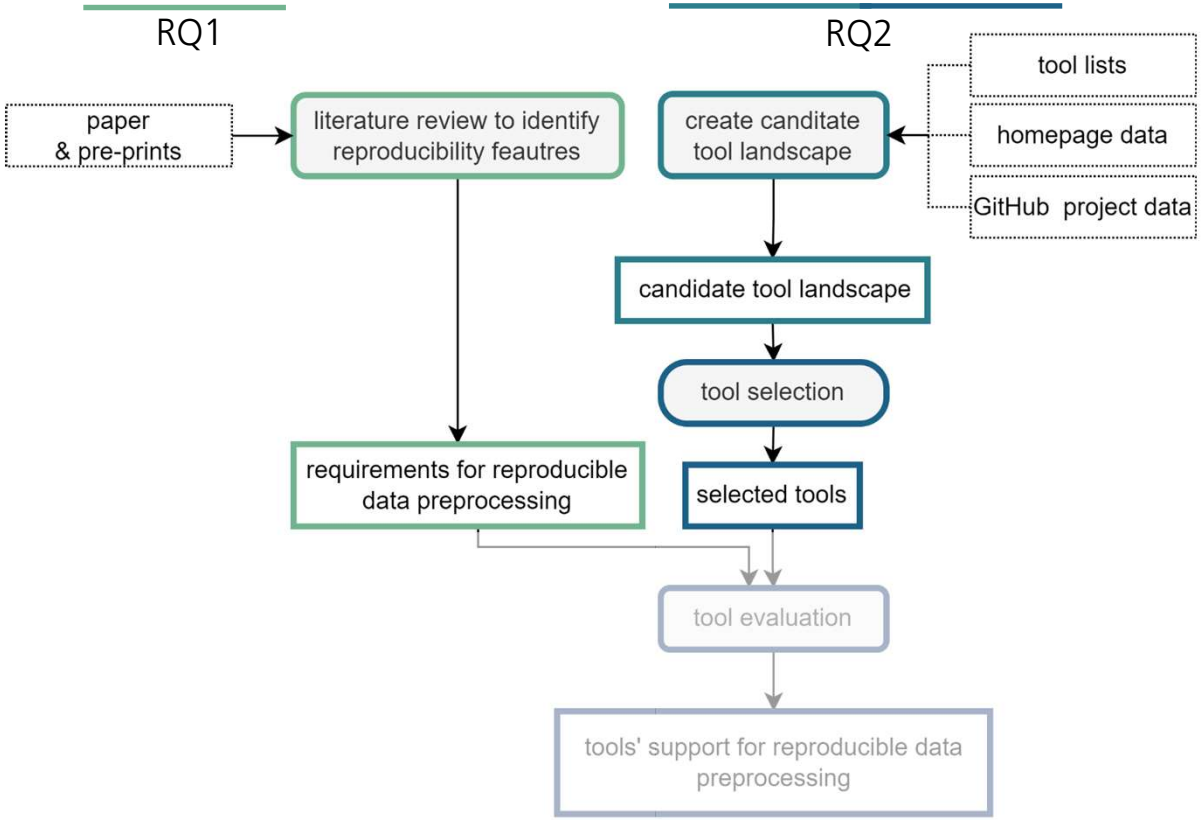

Methods and Design

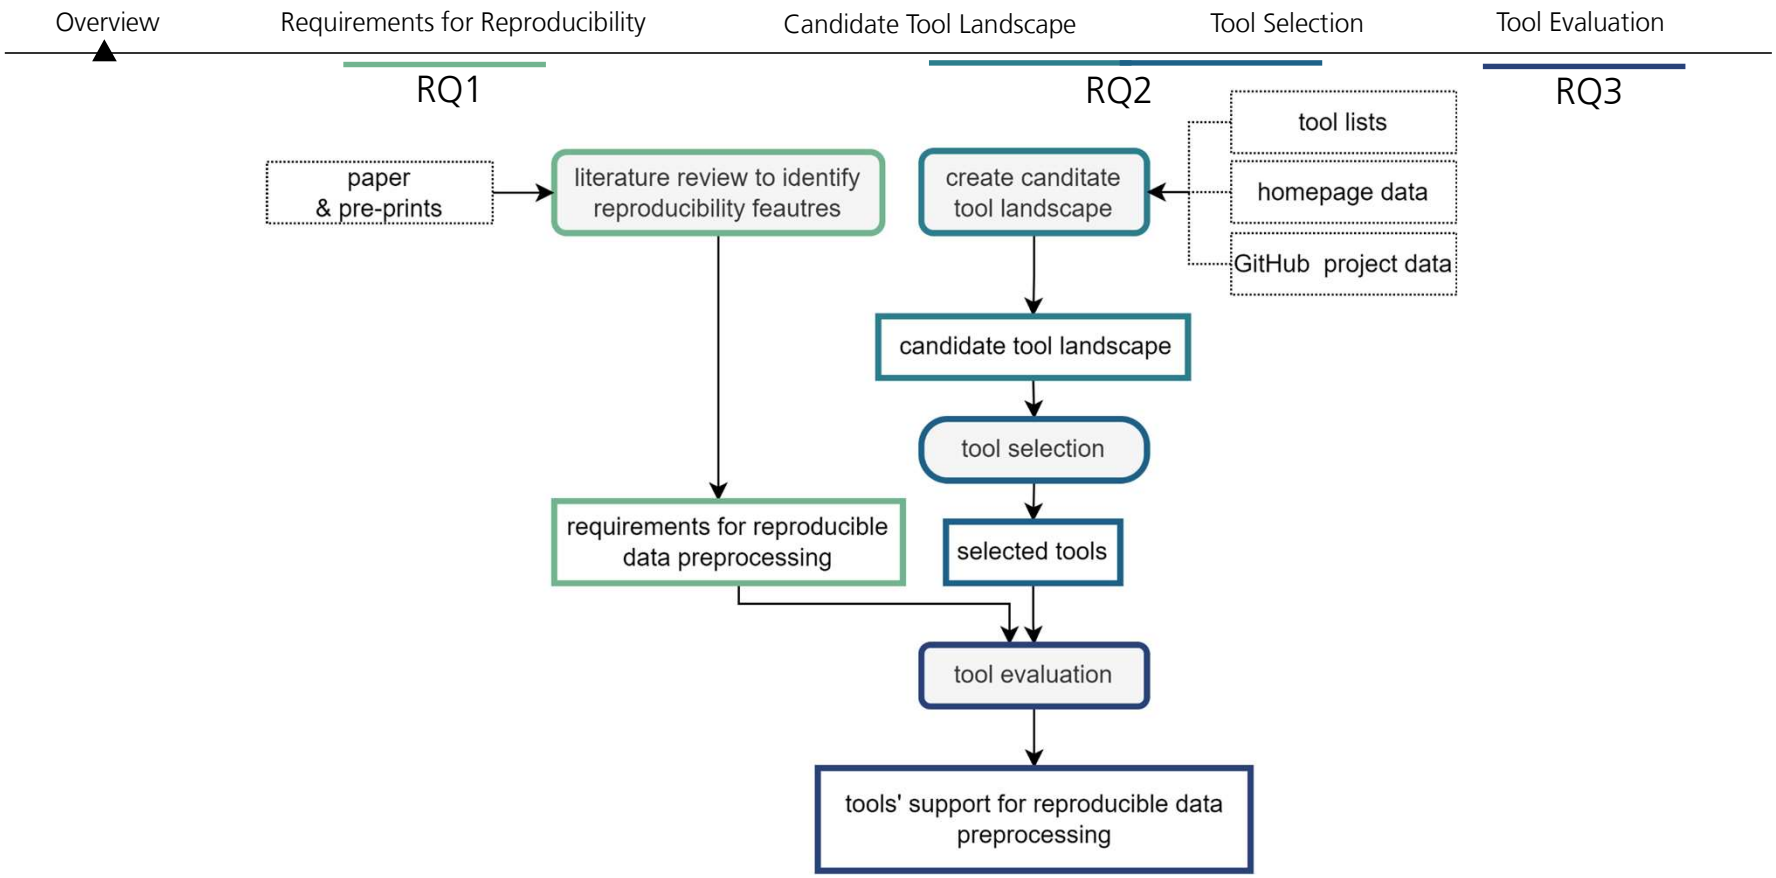

Methods and Design

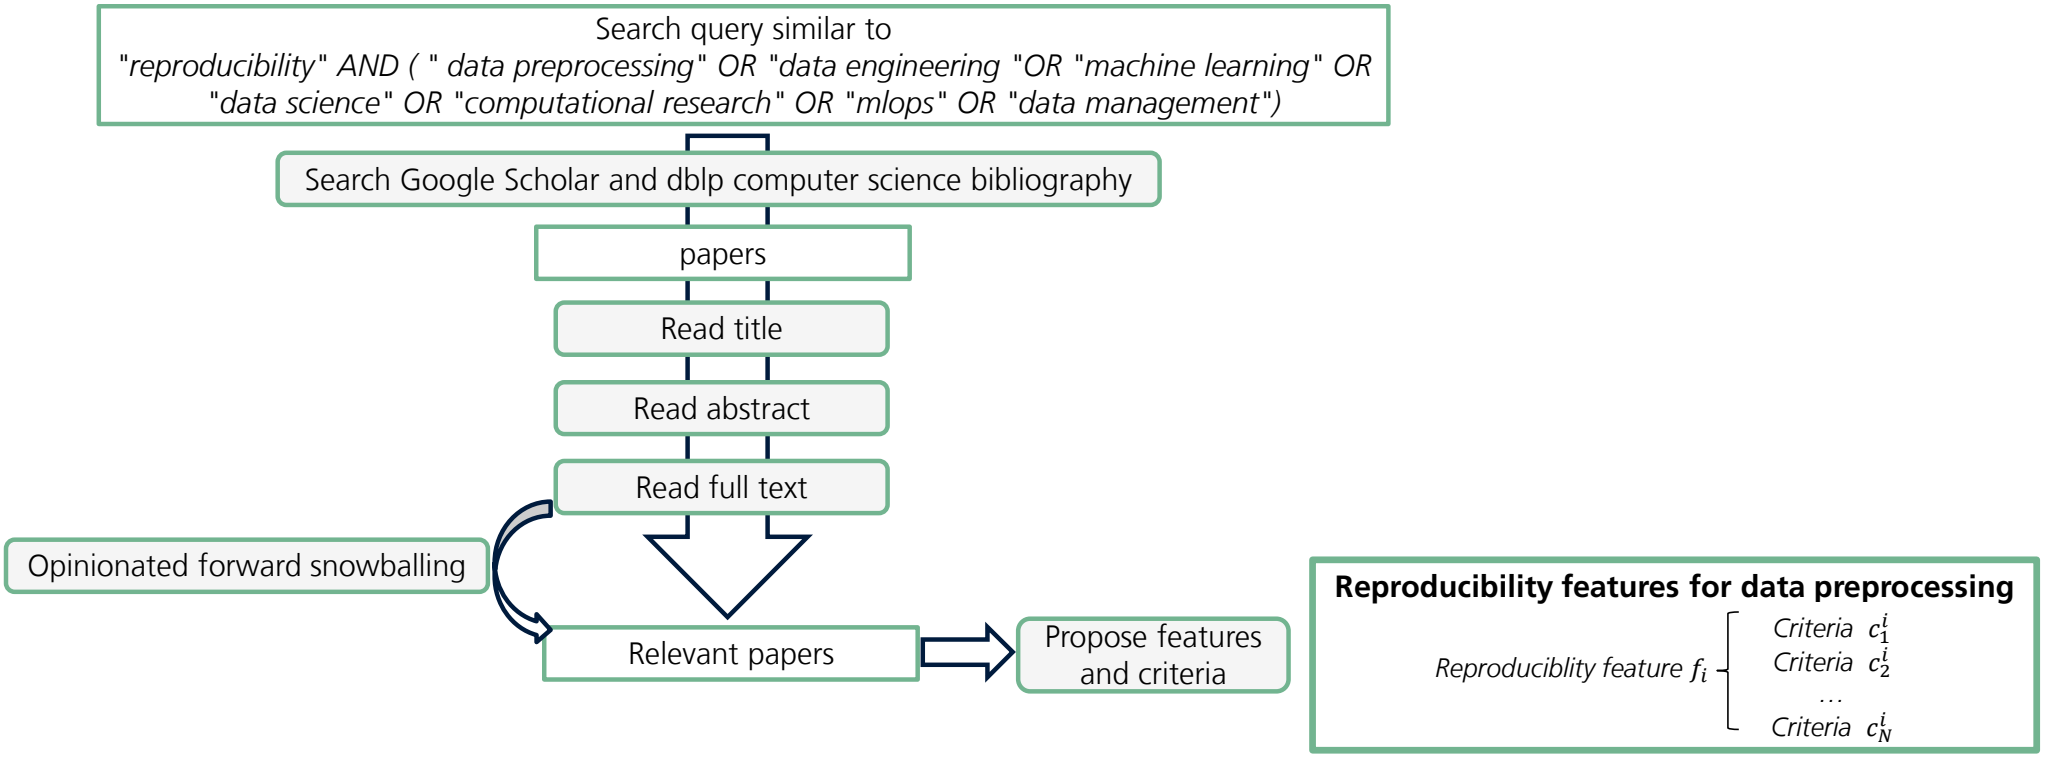

Methods and Design

Approach

Use community-maintained lists of tools in the AI and data domain to create a candidate tool landscape, and integrate GitHub and homepage data.

Methods and Design

Raw Data Sources

Existing Workflow Systems - https://s.apache.org/existing-workflow-systems

- incomplete list of computational analysis workflow systems
- Information per tool:
  - Tool name
  - Description (optional)
  - One or more Uniform Resource Locators (URLs) to the tool homepage, repository, or publication
- File format:
  - reStructuredText

Methods and Design

Raw Data Sources

Existing Workflow Systems - https://s.apache.org/existing-workflow-systems

1. Arvados - CWL-based distributed computing platform for data analysis on massive data sets. https://arvados.org/ https://github.com/arvados/arvados
2. Apache Taverna http://www.taverna.org.uk/ https://taverna.incubator.apache.org/
3. Galaxy http://galaxyproject.org/
4. SHIWA https://www.shiwa-workflow.eu/
5. Apache Oozie https://oozie.apache.org/
6. DNANexus https://wiki.dnanexus.com/API-Specification-v1.0.0/IO-and-Run-Specifications https://wiki.dnanexus.com/API-Specification-v1.0.0/Workflows-and-Analyses

Methods and Design

Raw Data Sources

Linux Foundation AI and Data Landscape- https://landscape.lfai.foundation/

- Interactive tool overview in the AI and data domain.
- Landscape is dynamically generated based on a YAML file in the corresponding GitHub repository
- Relevant information per tool:
  - Tool name
  - Category and Subcategory
  - Homepage URL
  - Repository URL
- File format:
  - YAML

Methods and Design

Raw Data Sources

Linux Foundation AI and Data Landscape- https://landscape.lfai.foundation/

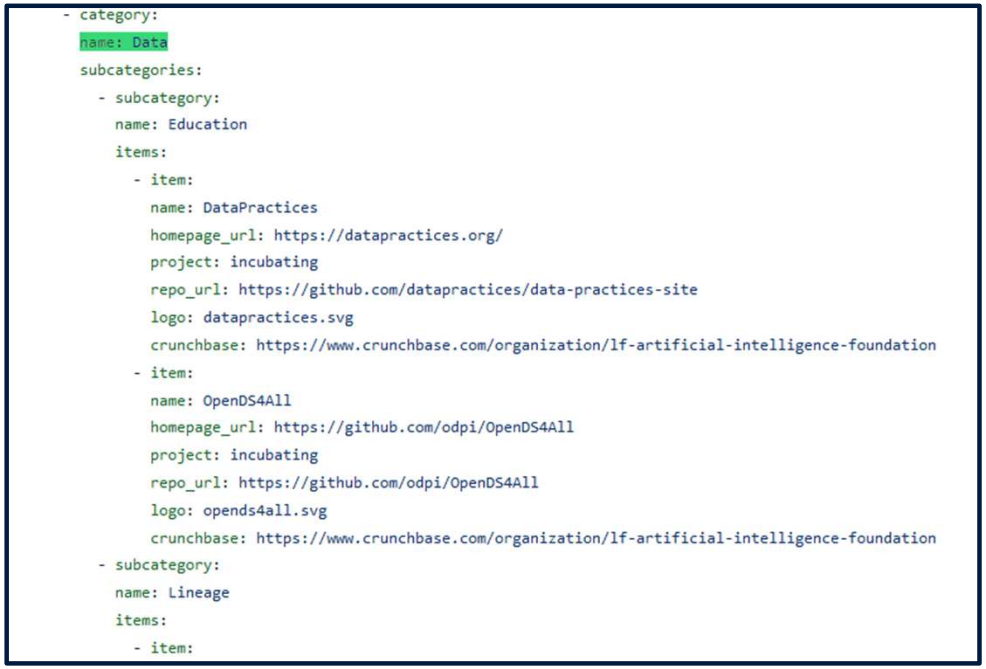

```
- category:
  name: Data
  subcategories:
    - subcategory:
      name: Education
      items:
        - item:
          name: DataPractices
          homepage_url: https://datapractices.org/
          project: incubating
          repo_url: https://github.com/datapractices/data-practices-site
          logo: datapractices.svg
          crunchbase: https://www.crunchbase.com/organization/lf-artificial-intelligence-foundation
        - item:
          name: OpenDS4All
          homepage_url: https://github.com/odpi/OpenDS4All
          project: incubating
          repo_url: https://github.com/odpi/OpenDS4All
          logo: opends4all.svg
          crunchbase: https://www.crunchbase.com/organization/lf-artificial-intelligence-foundation
    - subcategory:
      name: Lineage
      items:
        - item:
```

Methods and Design

Raw Data Sources

Awesome Pipeline- https://github.com/pditommaso/awesome-pipeline

- Community-curated list focusing on pipeline toolkits
- Information per tool:
  - Tool name
  - Short description
  - Grouped via sections and subsections
  - URL – either repository or homepage
- File format:
  - Markdown

Methods and Design

Raw Data Sources

Awesome Pipeline- https://github.com/pditommaso/awesome-pipeline

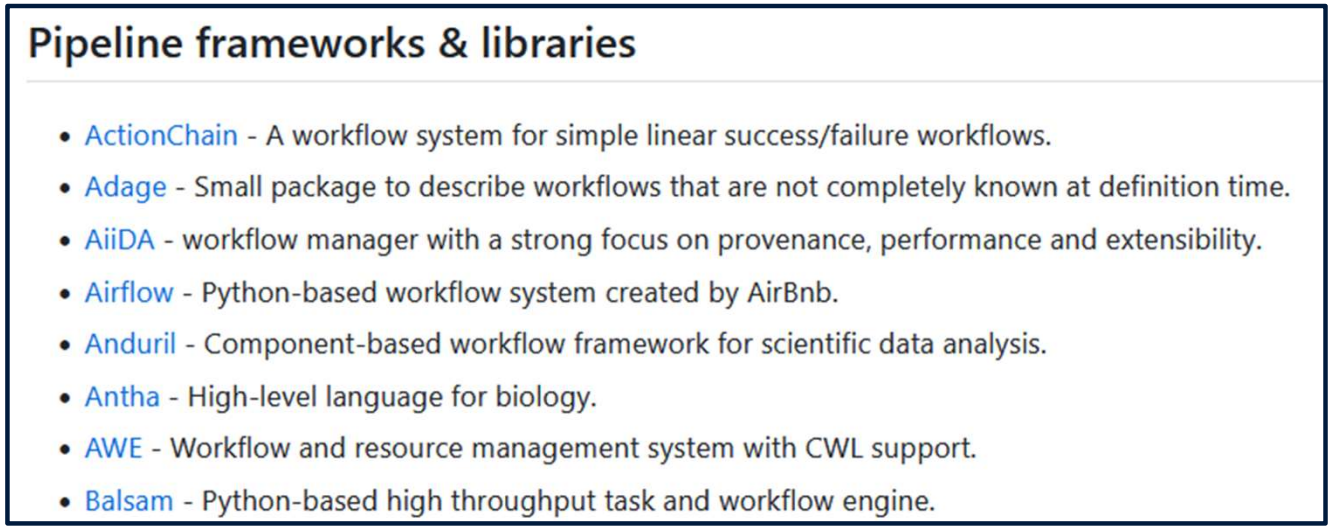

Pipeline frameworks & libraries

- ActionChain - A workflow system for simple linear success/failure workflows.
- Adage - Small package to describe workflows that are not completely known at definition time.
- AiiDA - workflow manager with a strong focus on provenance, performance and extensibility.
- Airflow - Python-based workflow system created by AirBnb.
- Anduril - Component-based workflow framework for scientific data analysis.
- Antha - High-level language for biology.
- AWE - Workflow and resource management system with CWL support.
- Balsam - Python-based high throughput task and workflow engine.

Methods and Design

Raw Data Sources

Awesome Data Engineering- https://github.com/igorbarinov/awesome-data-engineering

- Community-curated list focusing on data engineering tools
- Information per tool:
  - Tool name
  - Short description
  - Grouped via sections and subsections
  - URL – either repository or homepage
- File format:
  - Markdown

Methods and Design

Raw Data Sources

Awesome Data Engineering- https://github.com/igorbarinov/awesome-data-engineering

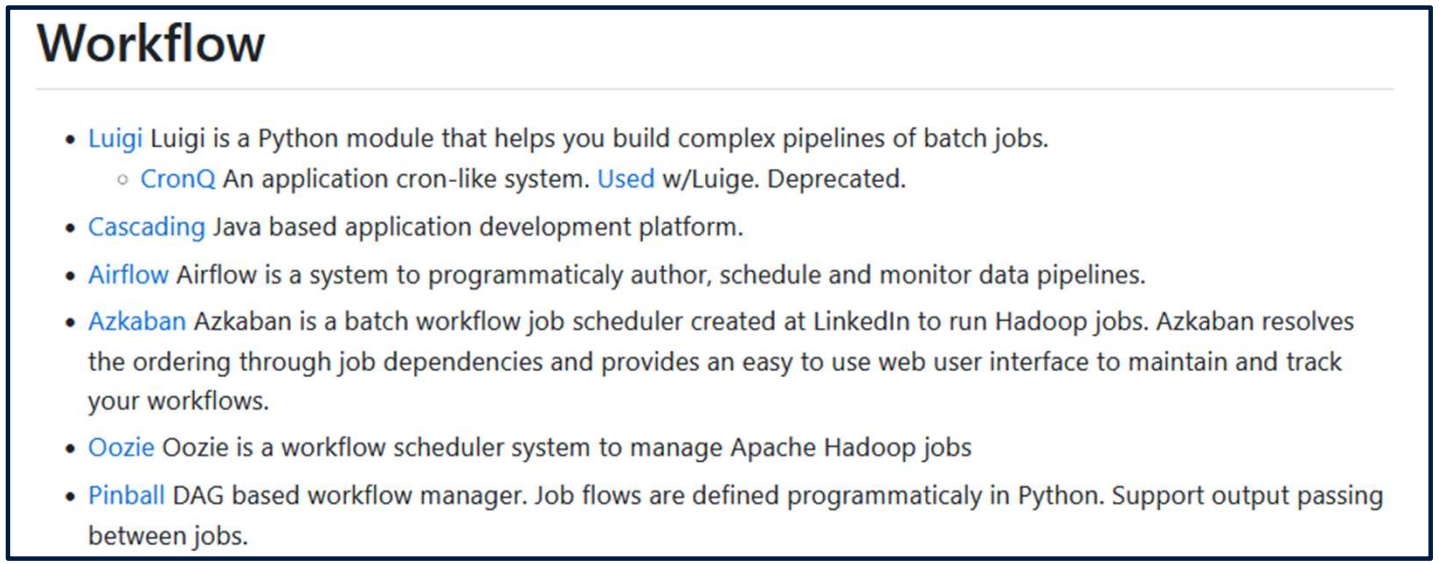

Workflow

- Luigi Luigi is a Python module that helps you build complex pipelines of batch jobs.
  - CronQ An application cron-like system. Used w/Luige. Deprecated.
- Cascading Java based application development platform.
- Airflow Airflow is a system to programmaticaly author, schedule and monitor data pipelines.
- Azkaban Azkaban is a batch workflow job scheduler created at LinkedIn to run Hadoop jobs. Azkaban resolves the ordering through job dependencies and provides an easy to use web user interface to maintain and track your workflows.
- Oozie Oozie is a workflow scheduler system to manage Apache Hadoop jobs
- Pinball DAG based workflow manager. Job flows are defined programmaticaly in Python. Support output passing between jobs.

Methods and Design

Data preprocessing pipeline to create the candidate tool landscape

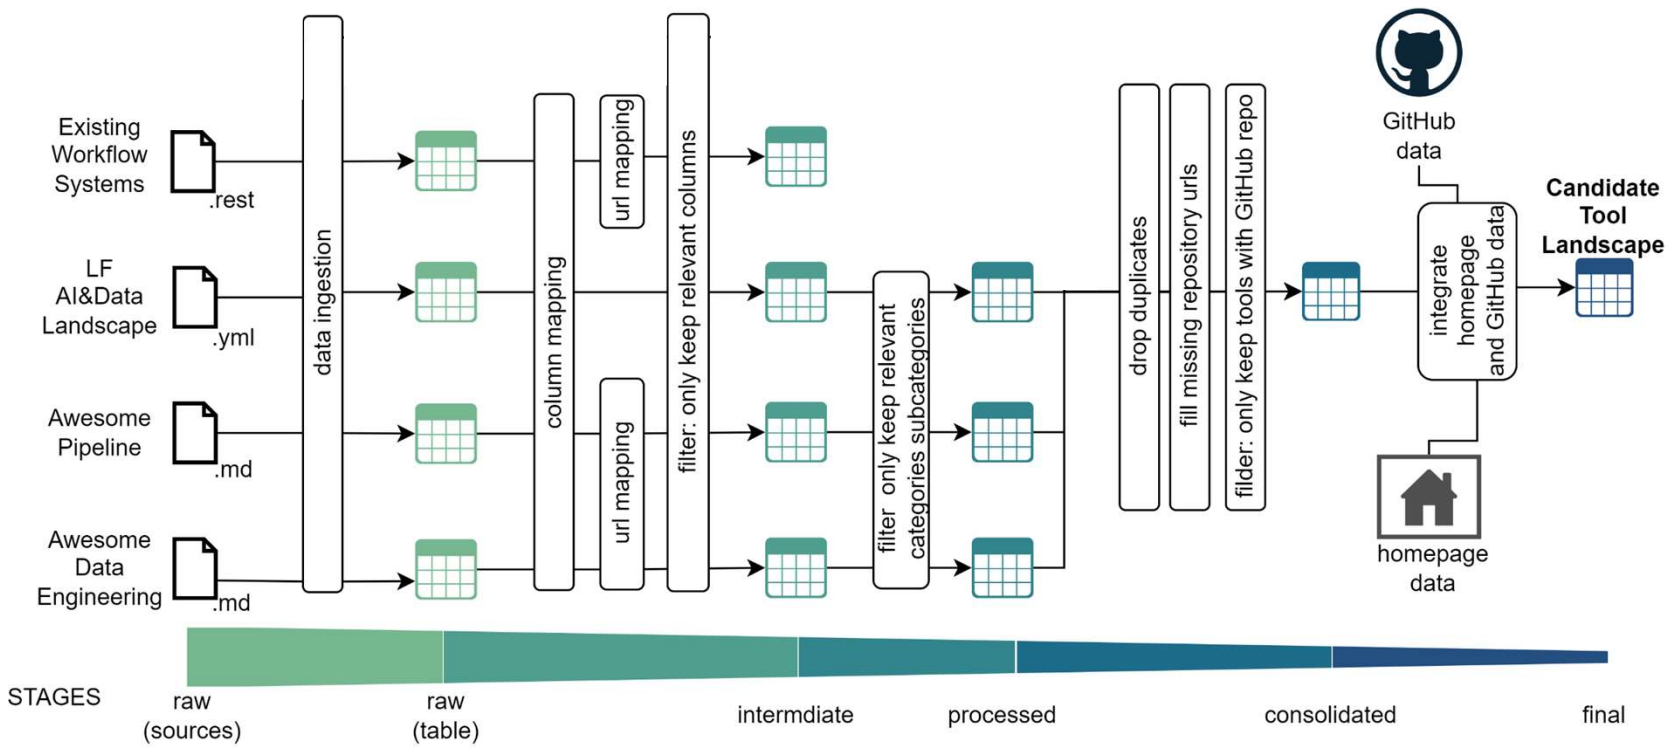

Methods and Design

Data preprocessing pipeline to create the candidate tool landscape

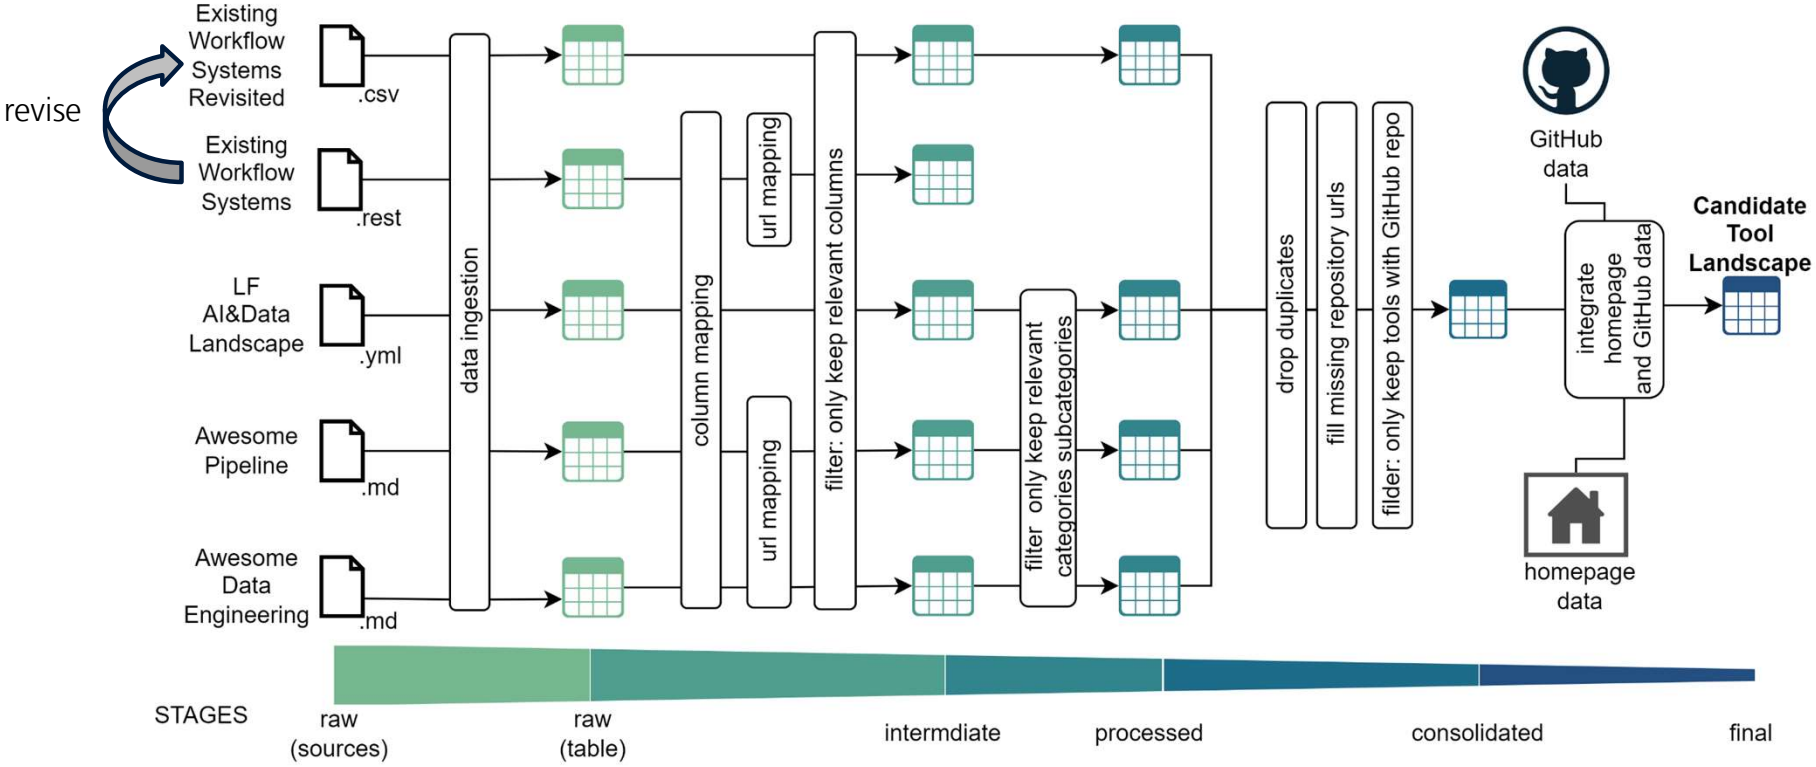

Methods and Design

- Conceptual Data Model

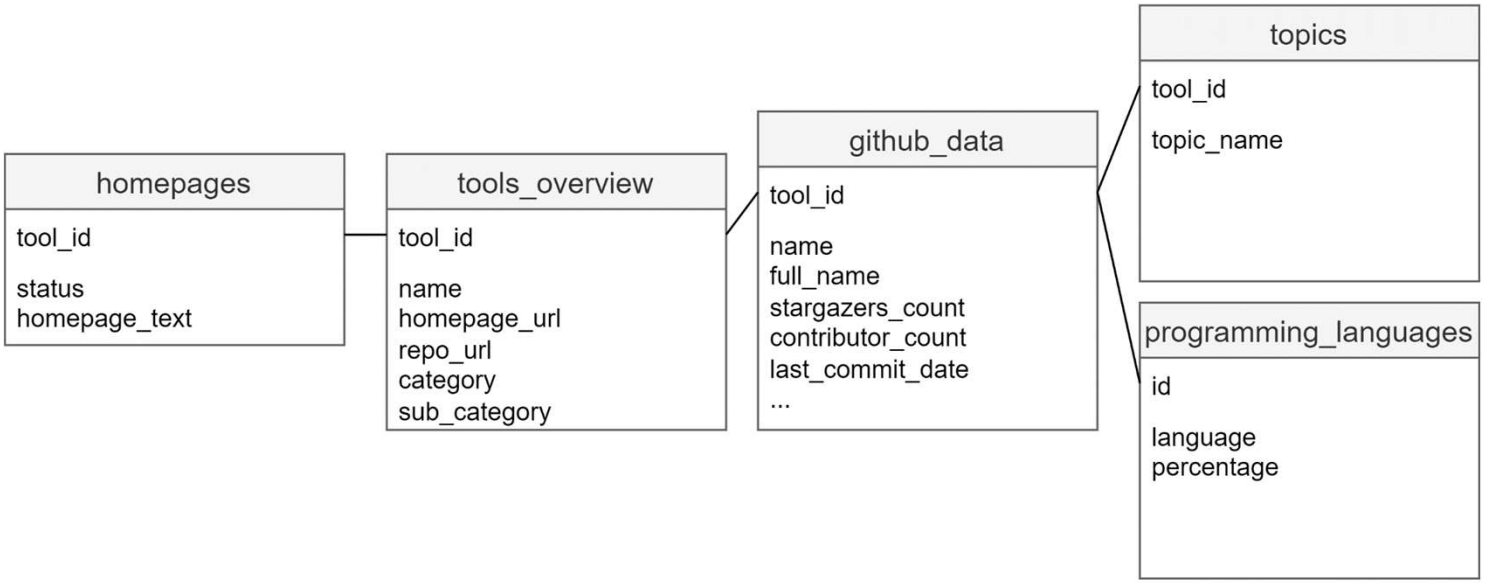

Methods and Design

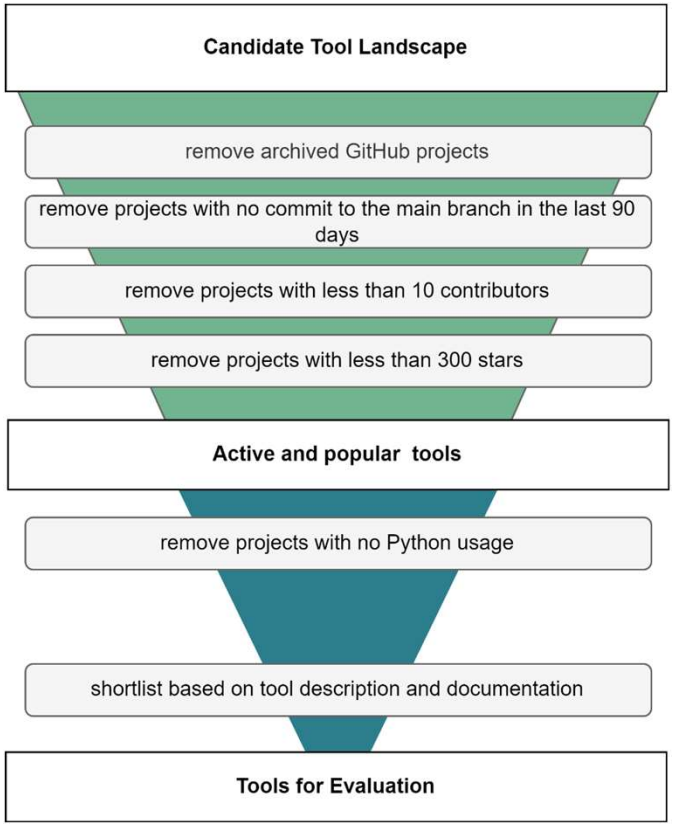

Methods and Design

1. Search the documentation and homepage for each selected tool to identify the support for each criterion of a reproducibility feature by assigning the support level.

Methods and Design

1. Search the documentation and homepage for each selected tool to identify the support for each criterion of a reproducibility feature by assigning the support level.

1.	unsupported	The tool does not support the criterion of a reproducibility feature.
2.	standard solution	The tool does not support the criterion of a reproducibility feature. However, this gap can be closed by a solution that the community sees as a default. For example, GitHub for public code hosting
3.	enterprise support	The tool supports the criterion of a reproducibility feature in the enterprise version, but the functionality is not available in the open-source version.
4.	integration	The tool proves an integration with a third-party solution, which supports the criterion of a reproducibility feature.
5.	partially	The tool partially supports the criterion of a reproducibility feature.
6.	full	The tool partially supports the criterion of a reproducibility feature.

Methods and Design

1. Search the documentation and homepage for each selected tool to identify the support for each criterion of a reproducibility feature by assigning the support level.
2. Quantify the support for a reproducibility feature

Methods and Design

1. Search the documentation and homepage for each selected tool to identify the support for each criterion of a reproducibility feature by assigning the support level.
2. Quantify the support for a reproducibility feature
    1. Assign a numeric value to each support level
        1. unsupported 0
        2. standard solution 0
        3. enterprise support 0
        4. integration 1
        5. partially 0
        6. full 1

Methods and Design

1. Search the documentation and homepage for each selected tool to identify the support for each criterion of a reproducibility feature by assigning the support level.
2. Quantify the support for a reproducibility feature
    1. Assign a numeric value to each support level
    2. Calculate mean of all criteria values for a feature

Methods and Design

1. Search the documentation and homepage for each selected tool to identify the support for each criterion of a reproducibility feature by assigning the support level.
2. Quantify the support for a reproducibility feature
    1. Assign a numeric value to each support level
    2. Calculate mean of all criteria values for a feature
3. Quantify the overall support for reproducible data preprocessing for each tool
    1. Calculate the mean of all reproducibility feature scores

Results

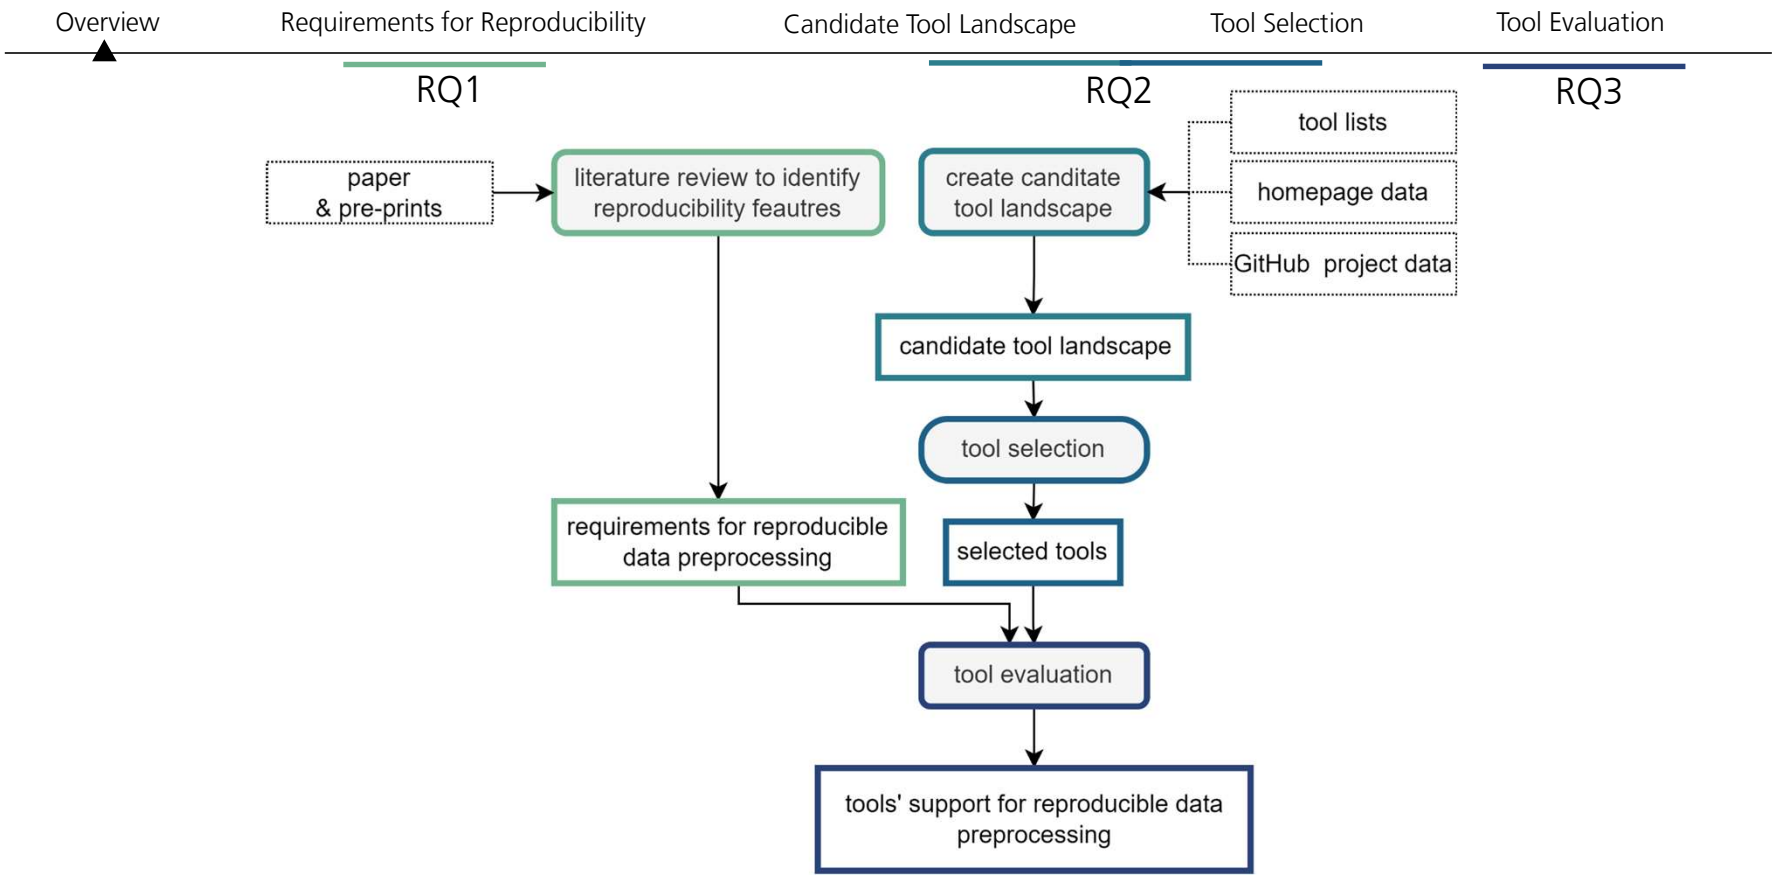
Overview
Requirements for Reproducibility
Candidate Tool Landscape
Tool Selection
Tool Evaluation
RQ1
RQ2
RQ3
paper
& pre-prints
literature review to identify
reproducibility feautres
create canditate
tool landscape
tool lists
homepage data
GitHub project data
candidate tool landscape
tool selection
requirements for reproducible
data preprocessing
selected tools
tool evaluation
tools' support for reproducible data
preprocessing

Results

Proposed reproducibility features for data preprocessing

Feature	Nr of criteria	Description
Code Sharing	3	Code is shared in a public code repository, versioned, and is citable.
Code Documentation	4	Code is documented and facilitated by a default structure and notebooks. User guides and static code analysis are supported.
Code License	2	A license is added to the project.
Code Review	2	A Code review process is described or integrated.
Workflow	4	Data preprocessing functions and configurations are abstracted in a workflow representation, such that the workflow is maintainable, portable, scalable, and documented.
Software and Code Dependencies	3	Software and code dependencies are specified in a standardized way using package manager and container.
Operating System	1	Operating System is specified as a part of a container image.
Kernel	1	A virtual machine image can be created.
Hardware	3	The hardware requirements are documented or specified declarative or as infrastructure as code.
SWE best practices	2	Testing and continuous integration are supported.
Data Sharing	3	Data is in a cloud storage, public repository, and is citable.
Data Documentation	3	Data is described in a basic way, annotated, or using a meta data standard.
Data License	2	A license is added to the data.
Data Quality	3	Data quality gates and measures are supported, via statistics, typing, schemas, and advanced data quality assessments.
Data Provenance	5	Data provenance is captured code agnostic, based on workflow implementation. Metadata is captured and stored in metadata management. Analysis of provenance data is supported.
Data Versioning	3	Data is versioned throughout the data lifecycle.

Results

Proposed reproducibility features and criteria for data preprocessing

Feature	Criteria
Code Sharing	Repository Version control Citable
Code Documentation	Structure Notebook User guide Static code analysis
Code License	Added Enforced
Code Review	Process Integration
Workflow	Portable Scalable Maintainable Metadata
Software and Code Dependencies	Package managment Container Captured
Operating System	Container
Kernel	VM image
Hardware	Documented Hosted Service IaC

Feature	Criteria
Data Sharing	Cloud Storage Repository Citable
Data Documentation	Described Meta data Meta data standard
Data License	Stored Enforces
Data Quality	Statistics Typing/Schema Quality
Data Provenance	Code agnostic Implementable Metadata Metadata managment Analysis
Data Versioning	Storage agnostic Automation Abstraciton
SWE Best Practices	CI Testing

Results

Number of tools for each data source and data preprocessing stage

Stage	EWSR	LFADL	AP	ADE	Total
raw	335	428	205	185	1153
intermediate	335	428	205	185	1153
processed	263	57	165	45	530
consolidated	238	41	59	26	364
final	236	41	57	25	359

EWSR: Existing Workflow Systems Revisited
LFADL: Linux Foundation AI and Data Landscape
AP: Awesome Pipeline
ADE: Awesome Data Engineering

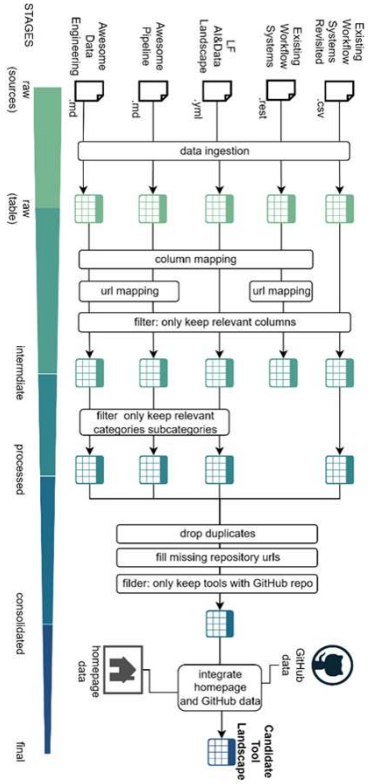

Results

Number of tools for each table where respective entries are valid and available

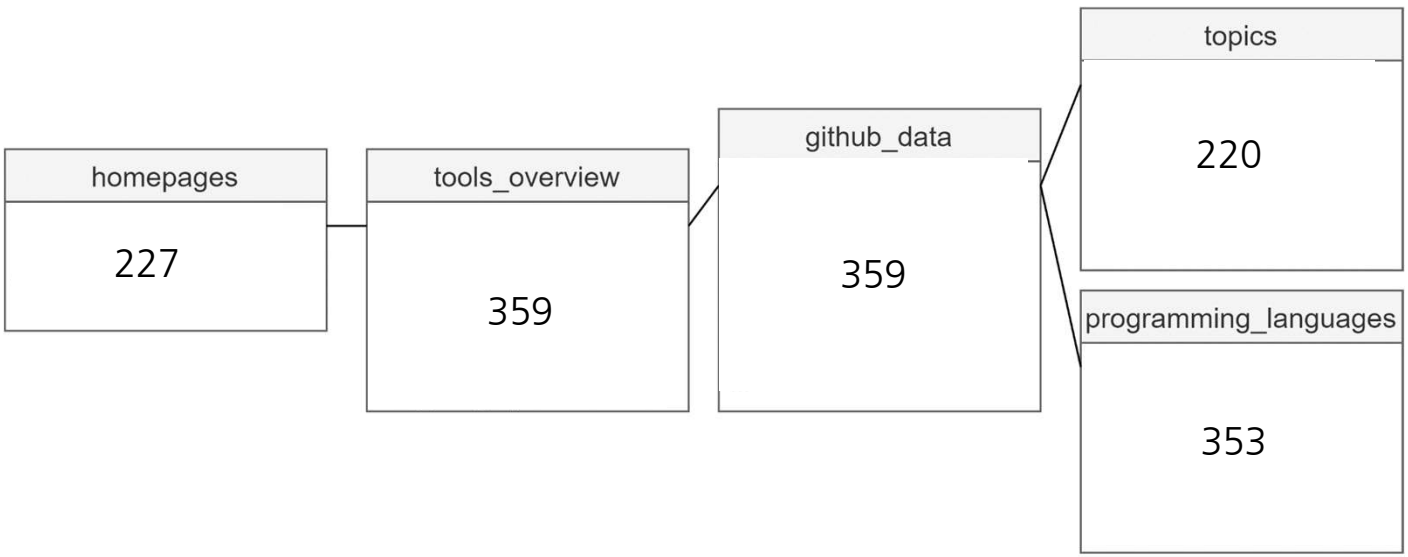

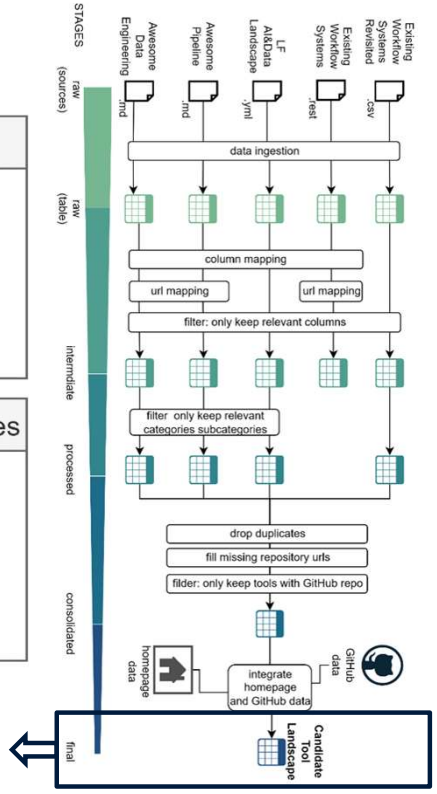

Results

Quality of the raw data sources with respect to
provided URLs for the homepage, repository, and publication.
The number of total and valid URLs (200 HTTP responses)
is given for the URL columns in the format: nr of URLS (nr of valid URLs).

nr. of	EWSR	EWS	LFADL	AP	ADE
tools	335	335	428	205	185
homepage URLs	252 (229)	209 (186)	428 (406)	89 (77)	106 (104)
publication URLs	199 (191)	75 (70)	0 (0)	0 (0)	0 (0)
repository URLs	263 (259)	158 (155)	340 (340)	116 (115)	79 (77)

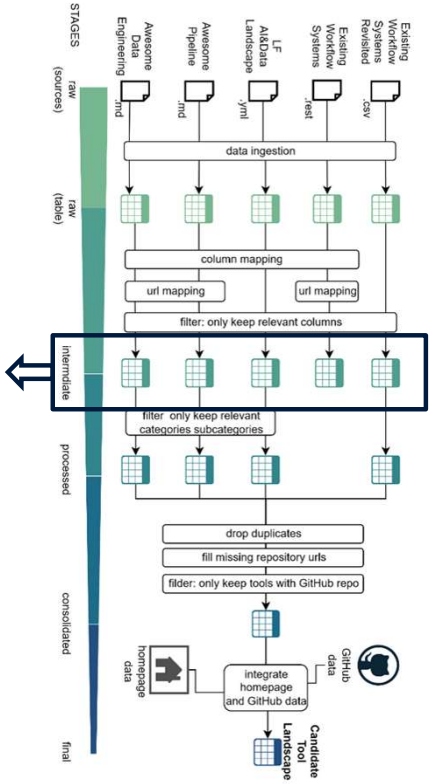

EWSR: Existing Workflow Systems Revisited
EWS: Existing Workflow Systems
LFADL: Linux Foundation AI and Data Landscape
AP: Awesome Pipeline
ADE: Awesome Data Engineering

Results

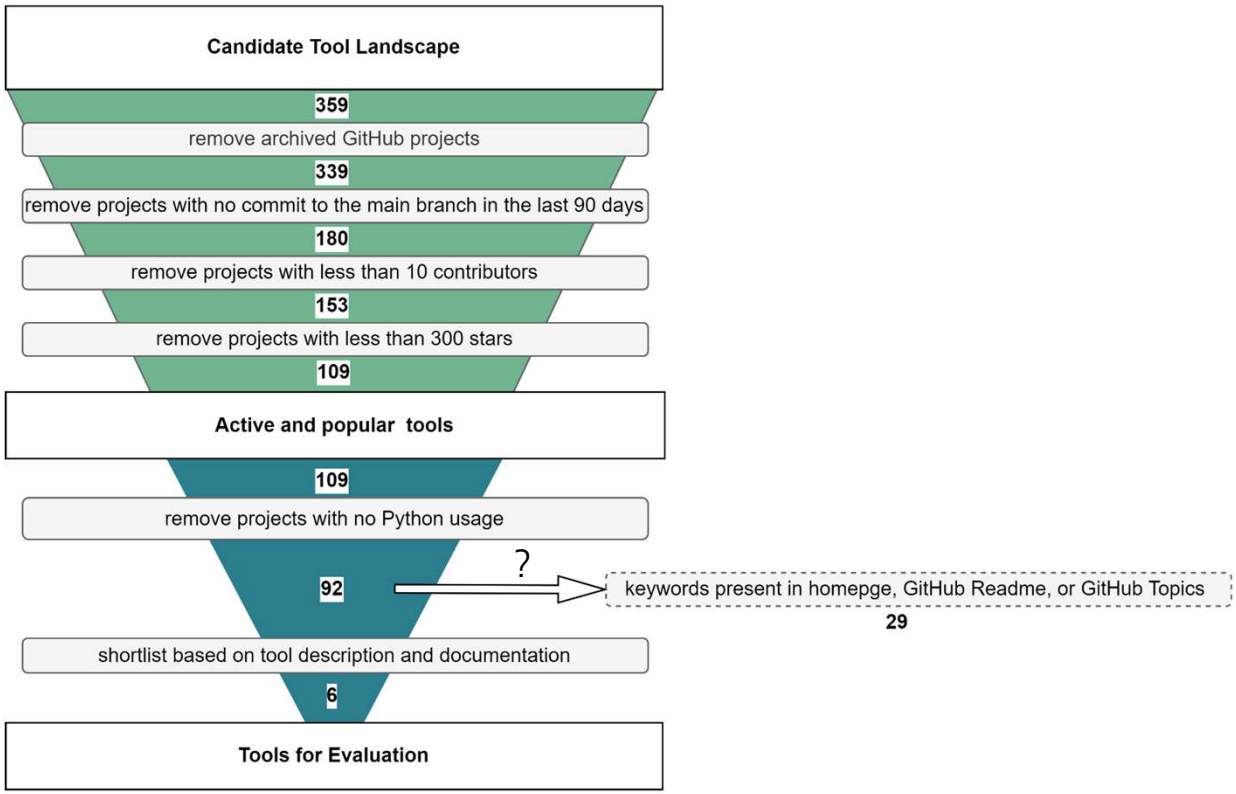
Candidate Tool Landscape
359
remove archived GitHub projects
339
remove projects with no commit to the main branch in the last 90 days
180
remove projects with less than 10 contributors
153
remove projects with less than 300 stars
109
Active and popular tools
109
remove projects with no Python usage
92
?
keywords present in homepge, GitHub Readme, or GitHub Topics
29
shortlist based on tool description and documentation
6
Tools for Evaluation

Results

Overview | Requirements for Reproducibility | Candidate Tool Landscape | Tool Selection | Tool Evaluation

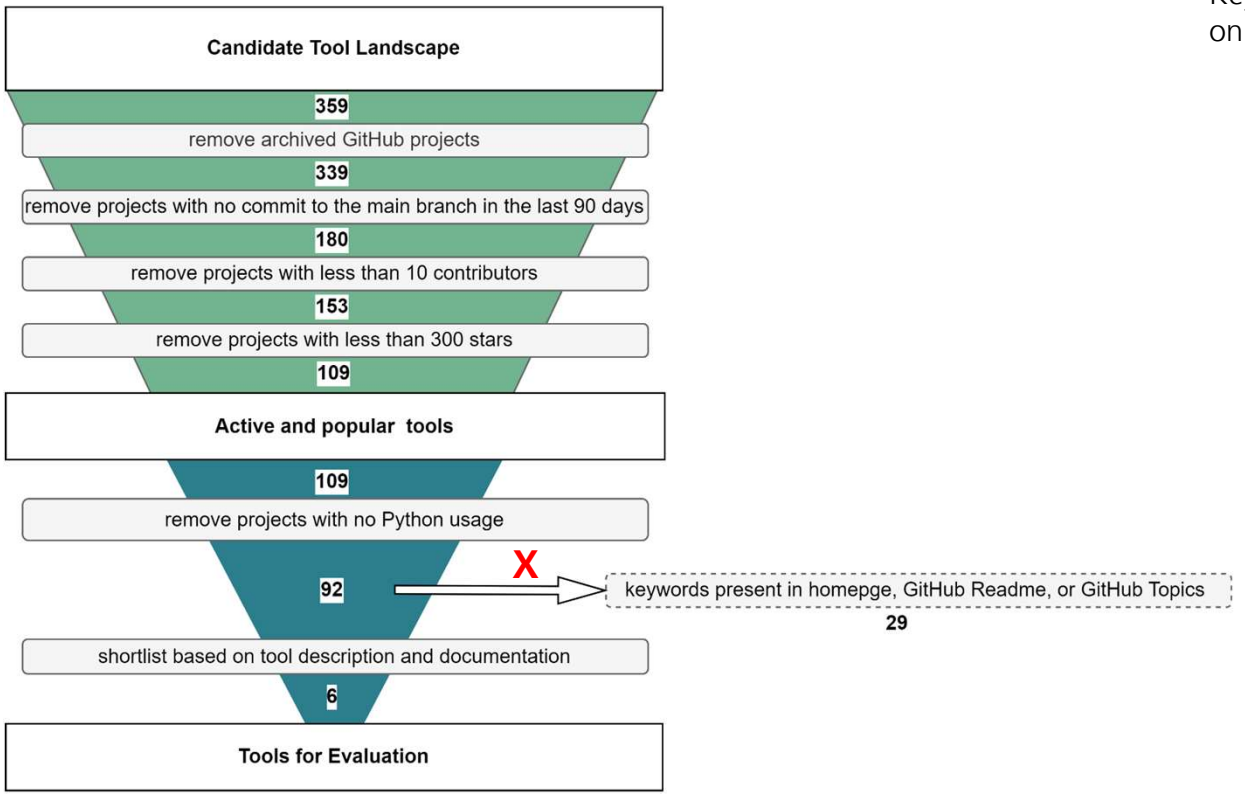

Keywords in the "reproducibility" word family occurred on the homepage, GitHub Readme, or GitHub Topics.

	Readme		Homepage		Topics	
	available	keyword	available	keyword	available	keyword
name						
mlflow	✓	✓	✓	✓	✓	-
dvc	✓	✓	✓	✓	✓	✓
kedro	✓	✓	✓	✓	✓	-
metaflow	✓	-	✓	-	✓	✓
enso	✓	✓	✓	✓	✓	-
pachyderm	✓	-	✓	✓	✓	-
BentoML	✓	-	✓	✓	✓	-
clearml	✓	✓	✓	✓	✓	-
kestra	✓	-	✓	✓	✓	-
flyte	✓	✓	✓	✓	✓	-
lakeFS	✓	✓	-	-	✓	-
joblib	✓	-	✓	✓	✓	-
zenml	✓	-	✓	✓	✓	-
Ax	✓	✓	✓	-	-	-
nextflow	✓	✓	✓	✓	✓	✓
snakemake	✓	✓	✓	✓	✓	✓
marquez	✓	-	✓	✓	✓	-
galaxy	✓	-	✓	✓	✓	-
cromwell	✓	✓	✓	-	✓	-
nipype	✓	✓	✓	✓	✓	-
wdl	✓	✓	✓	-	✓	✓
ck	✓	✓	✓	✓	✓	✓
covalent	✓	-	✓	✓	✓	-
datalad	✓	✓	✓	✓	✓	-
redun	✓	✓	✓	-	✓	-
jug	✓	✓	✓	✓	✓	-
aiida-core	✓	✓	✓	✓	✓	-
arvados	✓	✓	✓	✓	✓	-
cwltool	✓	-	✓	✓	✓	-

Results

Selected tools for evaluation with respect to reproducible data preprocessing

Tool	Category	Description	Stars	Contributors
Airflow	Data pipelines	A platform to programmatically author, schedule, and monitor workflows.	30866	418
Prefect	Data pipelines	Prefect is a workflow orchestration tool empowering developers to build, observe, and react to data pipelines.	12278	170
Dagster	Data pipelines	An orchestration platform for the development, production, and observation of data assets.	7840	293
dbt	Data warehouse transformation workflows	dbt enables data analysts and engineers to transform their data using the same practices that software engineers use to build applications.	7230	256
Flyte	Data pipelines	Scalable and flexible workflow orchestration platform that seamlessly unifies data, ML and analytics stacks.	3561	121
Snakemake	Bioinformatic workflows	The Snakemake workflow management system is a tool to create reproducible and scalable data analyses.	1749	267

Results

Reproducibility support for each criterion of a reproducibility feature with respect to the evaluated tools

Feature	Criteria	airflow	dagster	prefect	flyte	dbt	snakemake
Code Sharing	Repository	standard solution	enterprise support	standard solution	standard solution	enterprise support	standard solution
	Version control	standard solution	enterprise support	standard solution	standard solution	enterprise support	standard solution
	Citable	standard solution	standard solution	standard solution	standard solution	standard solution	integration
Code Documentation	Structure	unsupported	full	partially	full	full	partially
	Notebook	integration	integration	integration	integration	integration	integration
	User guide	standard solution	standard solution	standard solution	partially	full	partially
	Static code analysis	integration	standard solution	standard solution	standard solution	partially	full
Code License	Added	standard solution	standard solution	standard solution	partially	standard solution	partially
	Enforced	unsupported	unsupported	unsupported	unsupported	unsupported	unsupported
Code Review	Process	standard solution	enterprise support	standard solution	standard solution	standard solution	standard solution
	Integration	standard solution	enterprise support	standard solution	standard solution	integration	standard solution
Workflow	Portable	full	full	full	full	partially	full
	Scalable	full	full	full	full	full	full
	Maintainable	integration	full	full	full	full	full
	Metadata	partially	full	full	full	full	unsupported
Software and Code Dependencies	Package managment	full	full	full	full	full	full
	Container	full	full	full	full	integration	full
	Captured	unsupported	unsupported	unsupported	unsupported	unsupported	full
Operating System	Container	full	full	full	full	integration	full
Kernel	VM image	standard solution	standard solution	standard solution	standard solution	standard solution	standard solution
Hardware	Documented	standard solution	standard solution	partially	partially	standard solution	standard solution
	Hosted Service	full	full	full	full	unsupported	standard solution
	IaC	full	full	full	partially	standard solution	partially

Feature	Criteria	airflow	dagster	prefect	flyte	dbt	snakemake
Data Sharing	Cloud Storage	full	full	full	full	full	full
	Repository	standard solution	standard solution	standard solution	standard solution	standard solution	standard solution
	Citable	standard solution	standard solution	standard solution	standard solution	standard solution	integration
Data Documentation	Described	unsupported	full	partially	partially	full	standard solution
	Meta data	standard solution	full	full	partially	full	standard solution
	Meta data standard	integration	integration	integration	standard solution	standard solution	standard solution
Data License	Stored	standard solution	standard solution	standard solution	standard solution	standard solution	standard solution
	Enforces	unsupported	unsupported	unsupported	unsupported	unsupported	unsupported
Data Quality	Statistics	standard solution	full	full	full	enterprise support	standard solution
	Typing/Schema	standard solution	standard solution	standard solution	full	full	partially
	Quality	integration	integration	integration	integration	full	standard solution
Data Provenance	Code agnostic	unsupported	unsupported	unsupported	unsupported	unsupported	partially
	Implementable	partially	full	partially	full	full	standard solution
	Metadata	integration	integration	full	standard solution	integration	standard solution
	Metadata managment	integration	full	full	full	enterprise support	standard solution
	Analysis	partially	full	full	full	full	partially
Data Versioning	Storage agnostic	unsupported	full	full	full	full	unsupported
	Automation	unsupported	full	full	full	full	full
	Abstraciton	unsupported	full	full	full	full	unsupported
SWE Best Practices	CI	standard solution	enterprise support	standard solution	integration	enterprise support	integration
	Testing	partially	partially	partially	partially	integration	partially

unsupported
standard solution
enterprise support
integration
partially
full

Results

Reproducibility metric for each feature and tool

Feature	airflow	dagster	prefect	flyte	dbt	snakemake
Code Sharing	0.0	0.0	0.0	0.0	0.0	0.3
Code Documentation	0.5	0.5	0.5	0.8	1.0	1.0
Code License	0.0	0.0	0.0	0.5	0.0	0.5
Code Review	0.0	0.0	0.0	0.0	0.0	0.0
Workflow	1.0	1.0	1.0	1.0	1.0	0.8
Software and Code Dependencies	0.7	0.7	0.7	0.7	0.3	1.0
OS	1.0	1.0	1.0	1.0	0.0	1.0
Kernel	0.0	0.0	0.0	0.0	0.0	0.0
Hardware	0.7	0.7	1.0	1.0	0.3	0.3
Data Sharing	0.3	0.3	0.3	0.3	0.3	0.7
Data Documentation	0.3	1.0	1.0	0.7	0.7	0.0
Data License	0.0	0.0	0.0	0.0	0.0	0.0
Data Quality	0.3	0.7	0.7	1.0	0.7	0.3
Data Provenance	0.8	0.8	0.8	0.6	0.6	0.4
Data Versioning	0.0	1.0	1.0	1.0	1.0	0.0
SWE Best Practices	0.5	0.5	0.5	0.5	0.0	1.0
Tool Reproducibility Support	0.4	0.5	0.5	0.6	0.4	0.5

- $0.0 \leq s_f < 0.3$
- $0.3 \leq s_f \leq 0.7$
- $0.7 < s_f \leq 1.0$

Summary and Outlook

Summary | Outlook

- 16 features and respective criteria were proposed to help categorize the support for reproducible data preprocessing
- A candidate tool landscape in the AI and data domain was created to help identify relevant open-source tools in an iterative selection process
- A evaluation framework was designed, and six open-source tools were analyzed concerning their support for reproducible data preprocessing. None of them provides out-of-the-boc reproducibility.

Summary and Outlook

Summary | Outlook

- Further formalize the reproducibility features and criteria to facilitate identifying the support level by a tool
- Describe the default solution for a specific feature to indicate that it has not to be reinvented by a tool
- Evaluate the available integrations
- Outline a tool stack, which could further help to support reproducible data preprocessing

Q&A

Simon Grimm

References

- [Fei+20] Melanie Feinberg, Will Sutherland, Sarah Beth Nelson, Mohammad Hossein Jarrahi, and Arcot Rajasekar. "The New Reality of Reproducibility: The Role of Data Work in Scientific Research." In: Proceedings of the ACM on Human-Computer Interaction 4 (CSCW1 May 29, 2020)
- [Soh23] Emily Sohn. "The reproducibility issues that haunt health-careAI." In: Nature 613.7943 (Jan. 9, 2023)
- [AKV19] Stamatios-Aggelos N. Alexandropoulos, Sotiris B. Kotsiantis, and Michael N. Vrahatis. "Data preprocessing in predictive data mining." In: The Knowledge Engineering Review 34 (2019).
- [ATSO17] Mohammed Zuhair Al-Taie, Naomie Salim, and Adekunle Isiaka Obasa. "Successful Data Science Projects: Lessons Learned from Kaggle Competition." In: Kurdistan Journal of Applied Research 2.3 (Aug. 27, 2017).
- [FAM+97] A. Famili,Wei-Min Shen, RichardWeber, and Evangelos Simoudis. "Data Preprocessing and Intelligent Data Analysis." In: Intelligent Data Analysis 1.1 (Jan. 1, 1997)
- [GZ19] Carlos Vladimiro Gonzalez Zelaya. "Towards Explaining the Effects of Data Preprocessing on Machine Learning." In: 2019 IEEE 35th International Conference on Data Engineering (ICDE). 2019
- [BR21] Sumon Biswas and Hridesh Rajan. "Fair preprocessing: towards understanding compositional fairness of data transformers in machine learning pipeline." In: Proceedings of the 29th ACM Joint Meeting on European Software Engineering Conference and Symposium on the Foundations of Software Engineering. ESEC/FSE 2021.

References

- [IT18] Peter Ivie and Douglas Thain. "Reproducibility in Scientific Computing." In: ACM Computing Surveys 51.3 (July 16, 2018)
- [Mö+17] Steffen Möller et al. "Robust Cross-Platform Workflows: How Technical and Scientific Communities Collaborate to Develop, Test and Share Best Practices for Data Analysis." In: Data Science and Engineering 2.3 (Sept. 1, 2017)
- [Lei+21] Jeremy Leipzig, Daniel Nüst, Charles Tapley Hoyt, Karthik Ram, and Jane Greenberg. "The role of metadata in reproducible computational research." In: Patterns 2.9 (Sept. 10, 2021).
- [HO20] Matthew Hartley and Tjelvar S. G. Olsson. "dtoolAI: Reproducibility for Deep Learning." In: Patterns 1.5 (Aug. 14, 2020)
- [MJ21] Saúl Manzano and Adele C. M. Julier. "How FAIR are plant sciences in the twenty-first century? The pressing need for reproducibility in plant ecology and evolution." In: Proceedings of the Royal Society B: Biological Sciences 288.1944 (Feb. 10, 2021).
- [SLKR21] Sheeba Samuel, Frank Löffler, and Birgitta König-Ries. "Machine Learning Pipelines: Provenance, Reproducibility and FAIRData Principles." In: Provenance and Annotation of Data and Processes. Ed. by Boris Glavic, Vanessa Braganholo, and David Koop. Lecture Notes in Computer Science. Cham: Springer International Publishing, 2021
- [Wil+16] Mark D. Wilkinson et al. "The FAIR Guiding Principles for scientific data management and stewardship." In: Scientific Data 3.1 (Mar. 15, 2016)

References

- [Bar+22] Michelle Barker et al. "Introducing the FAIR Principles for research software." In: Scientific Data 9.1 (Oct. 14, 2022).
- [PRSA18] Beatriz Pérez, Julio Rubio, and Carlos Sáenz-Adán. "A systematic review of provenance systems." In: Knowledge and Information Systems 57.3 (Dec. 1, 2018)
- [Per19] Jeffrey M. Perkel. "Workflow systems turn raw data into scientific knowledge." In: Nature 573.7772 (Sept. 2, 2019).
- [MG22] Beatriz M. A. Matsui and Denise H. Goya. "MLOps: Five Steps to Guide its Effective Implementation." In: 2022 IEEE/ACM 1st International Conference on AI Engineering – Software Engineering for AI (CAIN). 2022
- [Pim+21] João Felipe Pimentel, Leonardo Murta, Vanessa Braganholo, and Juliana Freire. "Understanding and improving the quality and reproducibility of Jupyter notebooks." In: Empirical Software Engineering 26.4 (May 8, 2021)
- [Nü+20] Daniel Nüst, Vanessa Sochat, Ben Marwick, Stephen J. Eglen,Tim Head, Tony Hirst, and Benjamin D. Evans. "Ten simple rules for writing Dockerfiles for reproducible data science." In: PLOS Computational Biology 16.11 (Nov. 10, 2020).
- [Bar22